

OvO

OpenMP Versus Offload

Thomas Applencourt - apl@anl.gov

Argonne Leadership Computing Facility
Argonne National Laboratory
9700 S. Cass Ave
Argonne, IL 60349

Table of contents

1. Introduction
2. Mathematical Functions
3. Hierarchical Parallelism
4. Conclusion

Introduction

What is OvO?

- OpenMP versus Offload
- Not a conformance test suite!
- An OpenMP Offload **Mathematical** and **Hierarchical Parallelism** Test Suite

What is OvO? More technically

- 2700+ OpenMP 5.0 tests (C++ and FOR_{mula}TRAN_{slation})¹
- Scripts to compile, run, and check correctness
 - make
 - Bourne-Again SHell >3.2
 - Optionally monthlyPython3²
- *<https://github.com/TApplencourt/OvO>*

¹More precisely C++11 and F90 standards. Some math function are defined in C++17 and C++20 standards

²Used to generated summaries, or re-generate tests source-code

Why OvO?

1. OpenMP specification is **extensive**
2. Compilers only support a **subset** of it

For **application developers**:

- Check if a required feature is supported by a majority of compilers (else don't use it)

For **compiler developers**:

- Check if a required feature is supported by our compiler (else implement it)

How to Use OvO?

Help:

```
1  ./ovo.sh
2  Usage:
3    ovo.sh gen
4    ovo.sh run [<test_folder>...] [--no_long] [--no_loop]
5    ovo.sh display [--detailed | --failed | --passed ] [--no_long]
      ↪ [--no_loop] [<result_folder>...]
6    ovo.sh report [--no_long] [--no_loop] [<result_folder>]
7    ovo.sh clean
```

How to use OvO? “Live” Demo

```
1  $ git clone --quiet https://github.com/TApplencourt/OvO
2
3  $ CXXFLAGS='-std=c++2a' FC='gfortran' ./ovo.sh run
4  Running test_src/cpp/complex_cpp11 | Saving in ${long_path}
5  g++ -std=c++2a log_complex_float_complex_float.cpp -o
6  ↪ log_complex_float_complex_float.exe
7  [...]
8  timeout -k 5s 10s ./log_complex_float_complex_float.exe
9  [...]
10 $ ./ovo.sh report
11 >> test_result/2020-05-12_21-02_iris15.ftm.alcf.anl.gov
12 > cpp math and complex: 259 / 259 ( 100% ) pass
13   [failures: 0 compilation, 0 offload, 0 incorrect results]
14 > cpp hierarchical parallelism: 642 / 642 ( 100% ) pass
15   [failures: 0 compilation, 0 offload, 0 incorrect results]
16 > fortran math: 79 / 79 ( 100% ) pass
17   [failures: 0 compilation, 0 offload, 0 incorrect results]
18 > fortran hierarchical parallelism: 642 / 642 ( 100% ) pass
19   [failures: 0 compilation, 0 offload, 0 incorrect results]
20 > Summary: 1622 / 1622 ( 100% ) pass
21   [failures: 0 compilation, 0 offload, 0 incorrect results]
```


Mathematical Functions

How We Test?

- Floating-Point Arithmetic is hard. I'm Not an IEEE 754 expert
- Take a best-effort approach
- Compare GPU value to **assumed** gold CPU³
- Inputs are “goods” inputs:
 - No nan, no infinities, no subnormal, ...
 - Should not trigger any invalid operation, overflow, ...

³OpenMP specification doesn't define a required precision. OpenCL and CUDA® does. We choose to use a tolerance of 4 ulp, who correspond to Low Accuracy (LA) mode of Intel Math Kernel Library Vector Mathematics

What functions are we testing?

C++

cmath.h (391 tests for C++11, C++17, C++20), *complex.h* (51 tests) and one GNU extension (*sincos*)

FORTRAN

FORTRAN 77 Language Reference, section “Arithmetic and Mathematical Functions”⁴ (67 real and 17 complex tests)

⁴©1999 Sun Microsystems, Inc., 901 San Antonio Road, Palo Alto, California 94303-4900 U.S.A.

Mathematical Result (in no particular order)⁶

	Hardware	C++	FORTTRAN
Cray-llvm 9.1.3	V100	58%	100%
Cray-Classic 9.1.3	V100	49%	100%
clang-AOMP 11.0.1	Radeon VII	70%	44%
gcc 10.0.0	P100	5% ⁵	59%
clang 11.0.0	V100	70%	
xlC 16.1.1	V100	76%	100%

⁵I guess / hope it's a problem with our installation...

⁶--no_loop used

Hierarchical Parallelism

Goals

- One loop-nest
- All the combinations of OpenMP constructs:
 - team, distribute, parallel, for, loop, simd
- All the permutation of Type:
 - Real, **Complex**⁷
 - Single, Double precision
- Two big categories of tests:
 1. memory copy
 2. fold (reduction, atomic)

```
1 $find ./test_src/*/hierarchical_parallelism -type f | wc -l
2 2664
```

⁷OvO assume implicit mapping of `std::complex`. **Not part of the C++ OpenMP standard**

Quaternion of Extreme: One Line Complex Reduction

```
1  #pragma omp declare reduction(+: complex<float>: omp_out += omp_in)
2  std::complex<float> counter{};
3  #pragma omp target teams distribute parallel for simd
   ↪ reduction(+:counter) map(tofrom:counter)
4  for (int i = 0 ; i < N*M*L ; i++)
5      counter += std::complex<float> { 1.0f };
```

OpenMP features

- Custom reduction⁸
- In the combined statement, the *map* close and *reduction* close share a variable (OpenMP 5.0)

OpenMP “features”

Non-trivial type `complex<float>` is mapped

⁸Doesn't required in FORTRAN. Just saying that maybe one language is superior to another...

Quaternion of Extreme: exploded atomic

```
1  REAL counter = 0.
2  !$OMP TARGET map(tofrom:counter)
3  !$OMP TEAMS
4  !$OMP LOOP
5  DO i = 1 , L
6      !$OMP PARALLEL
7      !$OMP LOOP
8      DO j = 1 , N*M
9          !$OMP ATOMIC UPDATE
10         counter = counter + 1.
11     END DO
12     !$OMP END LOOP
13     !$OMP END PARALLEL
14 END DO
15 !$OMP END LOOP
16 !$OMP END TEAMS
17 !$OMP END TARGET
```

OpenMP Restriction

- No `!$OMP ATOMIC` in `!$OMP SIMD`
- No `!$OMP ATOMIC` with complex

Quaternion of Extreme: CUDA[®]-like reduction+atomic

```
1  double counter{};
2  #pragma omp target map(tofrom:counter)
3  #pragma omp teams
4  {
5      const int num_teams = omp_get_num_teams();
6      double partial_counter{};
7      #pragma omp parallel reduction(+: partial_counter)
8      {
9          const int num_threads = omp_get_num_threads();
10         partial_counter += double { 1.0f/(num_teams*num_threads) } ;
11     }
12     #pragma omp atomic update
13     counter += partial_counter;
14 }
```

Trivia: What is the final value of *counter*?

Quaternion of Extreme: host hierarchical parallelism

```
1  REAL counter = 0.
2  !$OMP PARALLEL DO REDUCTION(+: counter)
3  DO i = 1 , L
4      !$OMP TARGET TEAMS DISTRIBUTE PARALLEL DO MAP(TOFROM: counter)
5      DO j = 1 , N*M
6          !$OMP ATOMIC UPDATE
7          counter = counter + 1.
8      END DO
9      !$OMP END TARGET TEAMS DISTRIBUTE PARALLEL DO
10  END DO
11  !$OMP END PARALLEL DO
```

Hierarchical Parallelism Result (in no particular order)⁹

	Hardware	C++	FORTRAN
Cray-llvm 9.1.3	V100	70%	99%
Cray-Classic 9.1.3	V100	77%	99%
clang-AOMP 11.0.1	Radeon VII	50%	47%
gcc 10.0.0	P100	62%	81%
clang 11.0.0	V100	85%	
xlC 16.1.1	V100	73%	100%

⁹ `--no_loop` used

Conclusion

Conclusion

- OvO (<https://github.com/TApplencourt/OvO>)
- Nobody is perfect
- but FORTRAN is more perfect than others ¹⁰

¹⁰Cray and xLC have near a 100% pass rate. And FORTRAN have native complex support